

Drive the Thoughts: Runtime Monitoring of VLA Reasoning-Trajectory Consistency

Tian Yu, Lu Feng, Sebastian Elbaum
University of Virginia
Charlottesville, Virginia, USA
{buc9hh,lu.feng,selbaum}@virginia.edu

Abstract—Autonomous vehicles (AVs) operate in complex environments where failures are consequential. Sophisticated machine learning models for perception and planning are key to overcoming at least part of that complexity, but their black-box nature complicates validation and verification (V&V). The recent integration of Vision-Language-Action (VLA) models into AVs introduces a unique opportunity: besides generating trajectories, these models produce an explicit Chain-of-Thought (CoT) explaining their underlying rationale. This CoT provides a rich specification to cross-check model outputs and detect inconsistencies that may expose unsafe or unintended behavior. This paper assesses whether CoTs from a recent open driving VLA can support such monitoring. We curate DriveAlignBench, a specialized dataset from NVIDIA’s Alpamayo 1.5 VLA for AVs containing 150 CoT-trajectory pairs, which we manually annotate for reliability, trajectory consistency, and safety. Our analysis reveals that 33.3% of CoTs are unreliable. Among reliable CoTs, the generated trajectory is consistent with the CoT in 74% of cases. Leveraging this potential, we propose integrating a CoT-trajectory consistency check into a runtime monitor. The check is nontrivial: CoTs express open-vocabulary, scene-relative driving commitments, while trajectories are low-level ego-motion sequences whose semantics depend on road geometry and motion context. To bridge this gap, we develop a family of automated consistency monitors. Our best monitor, lane-relative F-LLM with GPT-5.5, achieves $F_1 = 0.75$, improving over the strongest raw-waypoint LLM baseline by +0.13 absolute F_1 and over a rule-based monitor by +0.38. We release DriveAlignBench, the monitor implementations, and annotation tools at <https://github.com/776styjsu/drive-the-thoughts>.

Index Terms—Runtime monitoring, autonomous vehicles, vision-language-action models, chain-of-thought reasoning, reasoning-action consistency, LLM-as-judge, software testing and verification

I. INTRODUCTION

Autonomous vehicles (AVs) must operate safely within an unpredictable physical world, where a single failure can lead to catastrophic consequences, as illustrated by the Uber automated-test-vehicle pedestrian fatality, the Tesla Autopilot Mountain View crash, and the Cruise driverless pedestrian-dragging incident [1]–[3]. Navigating this operational complexity is causing a paradigm shift in the underlying software architecture of AVs. Industry is moving from traditional, modular pipelines, where isolated machine-learning components handle narrow tasks such as perception, prediction, planning, or control, toward larger integrated subsystems and end-to-end multimodal driving models, including vision-language-action (VLA) models [4]–[9]. These frontier models ingest raw visual




Fig. 1. The two artifacts a reasoning VLA exposes at runtime, on one representative example. Given the (1) **scene** (the front-camera input to the model), the model emits a (2) **chain-of-thought** stating what it intends to do and why, together with a (3) **trajectory** τ that should realize the stated action. Here the CoT is *reliable*—grounded, safe, and coherent—and τ is *consistent* with it: the orange path shifts one lane left, exactly as the CoT asserts.

streams, navigation context, and telemetry sources to produce vehicle trajectories, and have shown promising performance on challenging and long-tail driving scenarios [5]–[7], [9].

Yet, this leap in capability magnifies a severe validation and verification (V&V) challenge because the multimodal billion-parameter space of modern models substantially expands the space of difficult-to-interpret failure modes. Concurrently, the black-box nature of deep neural networks makes rigorous V&V difficult: model internals are not expressed in the semantic vocabulary of formal driving specifications, such as vehicles, pedestrians, lanes, right-of-way, and temporal traffic interactions [10], [11]. When an end-to-end model misinterprets a complex intersection, miscalculates the intent of a merging vehicle, and generates an erroneous control action, V&V techniques struggle to assert that the behavior is incorrect given the high-dimensional input and the opacity of the model’s internal representations.

The recent integration of reasoning models into autonomous driving introduces a unique opportunity to address this challenge. Beyond merely generating raw trajectory coordinates, these models emit an explicit chain-of-thought-style reasoning (CoT) trace [6]–[9], [12], [13]. These thoughts are not generic language streams; in driving VLAs, they are intended to be grounded in the physical task of driving, capturing domain concepts such as lane geometry, traffic rules, agent interactions, and dynamic constraints. For example, the CoT in

Fig. 1, generated by NVIDIA’s Alpamayo 1.5 VLA [9], states a maneuver (change lane left) together with its scene-grounded justification (slower traffic ahead and a clear gap). From a V&V perspective, such CoT provides a rich specification that may serve as a semi-independent stream of evidence for cross-checking model outputs and flagging impending failures.

This use of internal reasoning for V&V rests on two assumptions that we systematically investigate in this paper:

- The generated CoT must be *reliable* (Fig. 1, panel 2).
- The CoT and proposed trajectories must be comparable for semantic *consistency* (Fig. 1, panels 2 and 3).

Our first contribution targets the first assumption. In Section III, we conduct an assessment of a recent open driving model trained to produce CoT together with trajectories. We curate DriveAlignBench, a specialized dataset from NVIDIA’s Alpamayo 1.5 VLA for autonomous vehicles [9], containing 150 CoT–trajectory pairs that we annotate for reliability, trajectory consistency, and safety. Our analysis reveals that reliability varies across driving-decision categories: among categories with at least 10 examples, reliability ranges from 30.0% to 81.3%, with 33.3% of CoT being unreliable overall. Among 100 reliable CoTs, reasoning aligns with actual trajectories in 74% of cases.

Our second contribution targets the second assumption by formulating CoT–trajectory agreement as a runtime monitoring problem. In Section IV, we ask whether the generated trajectory realizes the maneuver asserted by the CoT. This comparison is difficult because CoTs state open-vocabulary, scene-relative commitments, such as lane changes, yielding, and gap keeping, whereas trajectories expose only low-level ego motion whose interpretation depends on road geometry and scene context. We therefore design a family of automated monitors that compare CoTs and trajectories under increasingly rich trajectory representations. Our evaluation in Section V shows that adding lane-relative road geometry to an LLM judge yields the strongest monitor, F-LLM, with $F_1 = 0.75$, a +0.13 absolute improvement over the strongest raw-waypoint LLM baseline. The gain from this representation change is larger than the gain from switching to a stronger judge model, suggesting that scene-relative evidence is more important than judge scale for this task. We also find that inconsistency is safety-relevant: among reliable CoTs, human-labeled inconsistencies account for 7 of the 9 unsafe trajectories, and F-LLM monitors detect these 7 cases, unlike the raw-waypoint and rule-based baselines.

Reproducibility: We release DriveAlignBench, our monitoring architecture, and experiments at <https://github.com/776styjsu/drive-the-thoughts>.

II. BACKGROUND AND RELATED WORK

A. Background

Our work focuses on the analysis and monitoring of driving VLA models. More broadly, end-to-end driving systems commonly map sensor and context inputs to a planned future ego

trajectory [4]. In reasoning driving VLAs such as LINGO-2 and Alpamayo, the model can condition on visual observations together with route, language, or ego-state context, and can output both text reasoning or commentary and a future trajectory [5], [9]. Let $o_{\leq t}$ be the historical observations (including visual observations and historical ego-state) and ℓ be the natural-language driving prompt, which is fixed in our experiments. The VLA model studied here outputs:

$$(c_t, \tau_t) = f_\theta(o_{\leq t}, \ell)$$

where $c_t \in \mathcal{C}$ describes the intended driving behavior in text, and $\tau_t = (s_{t+1}^*, \dots, s_{t+H}^*) \in \mathcal{T}$ is the generated trajectory whose states carry waypoint positions and headings. Our proposal is to build monitors that check whether τ_t is consistent with c_t .

B. Reasoning in Driving VLAs

Building on chain-of-thought prompting in language models [14], recent driving VLA systems couple explicit reasoning with action generation to improve trajectory quality, efficiency, or robustness in complex scenarios. AdaThinkDrive [15] learns to switch between direct prediction and CoT reasoning based on scene complexity; AutoDrive-R² [16] uses a four-step CoT with self-reflection for supervised fine-tuning and GRPO with physics-grounded rewards; and Alpamayo models emit Chain-of-Causation traces describing driving decisions and causal factors [7], [9]. Alpamayo-R1 further uses RL post-training to improve reasoning quality and enforce reasoning–action consistency [7]. Across these systems, however, the CoT is mostly optimized as a means to improve trajectory quality but not as a runtime-checked artifact in its own right. This leaves two assumptions unverified at runtime: that the reasoning is trustworthy, and that the trajectory realizes it. Alpamayo-R1 targets the second during training through reasoning–action consistency rewards, but neither is checked as the vehicle drives.

C. CoT Faithfulness and Reasoning–Action Consistency

Prior work on chain-of-thought faithfulness cautions that stated reasoning can be a post-hoc rationalization rather than the model’s true decision process [17], [18]. In robotics, an observable form of this problem is inconsistency between the stated reasoning and the generated action. Closest to us, two efforts examine reasoning–action consistency directly. Mayumu et al. study Alpamayo-R1 reasoning faithfulness by measuring entity fidelity, action fidelity, reasoning–action consistency, and perturbation sensitivity using keyword extraction, ground-truth checks, and kinematic predicates [19]; Wu et al. improve tabletop manipulation robustness by sampling candidate action sequences at runtime, predicting their outcomes in simulation, and using a pretrained VLM to select an execution whose outcome aligns with the model’s textual plan [20]. We share their premise that reasoning–action consistency is checkable, but we study the challenges of using it for runtime monitoring in driving, where reasoning is often scene-relative while the generated trajectory is expressed as raw control




Fig. 2. Representative CoT-trajectory pairs from DriveAlignBench, each showing the front-camera view supplied to Alpamayo 1.5 and the CoT it emitted. (a,b) Reliable CoTs from very different scenes—dusk highway lead-following and a night approach to a red signal—are grounded, safe, and coherent. (c,d) Unreliable CoTs fail for different reasons: (c) a *grounding* failure that places a lead vehicle in the ego lane when the only nearby vehicle is in the adjacent lane, and (d) a *safety* failure that nudges into the opposing lane to overtake while an oncoming vehicle approaches. The decoded trajectory is withheld, since reliability is judged from the CoT and scene alone (Section III-A).

output. Our monitor bridges this representation gap rather than assuming that reasoning statements can be directly matched to trajectory predicates or candidate actions.

D. Runtime Monitoring and Verification of Driving Policies

Runtime monitoring offers a complementary way to assure learning-enabled driving systems when offline verification is insufficient. Prior work includes model-based safety envelopes, such as Responsibility-Sensitive Safety, which defines formal kinematic safety constraints for safe planned motion [21]; monitors for learned components, such as out-of-distribution detection and introspective failure prediction [22], [23]; and semantic runtime monitors that check authored driving properties over relational structures such as scene graphs and temporal scene-flow specifications [10], [11]. These approaches provide external safety or reliability checks over the planner, perception stack, or scene representation. Our work instead monitors whether the VLA’s generated trajectory is consistent with its own CoT description of the intended maneuver.

III. CoT RELIABILITY ANALYSIS

A CoT-trajectory consistency check is meaningful only if the CoT is itself a valid description of what the vehicle should do. Otherwise, the monitor output becomes difficult to interpret: a trajectory may be consistent with an ungrounded or unsafe rationale, or may be flagged as inconsistent simply because the CoT is not a reliable specification in the first place. We therefore first ask when a driving CoT can serve as a *candidate specification* for the monitors introduced in Section IV. We study this question on 150 CoT-trajectory pairs generated by Alpamayo 1.5 and address two research questions:

- **RQ1:** What proportion of CoTs are grounded in the scene, prescribe a safe action, and connect their evidence to that action coherently?
- **RQ2:** For which driving decisions and causal factors are CoTs most and least often reliable candidate specifications?

A. When Is a CoT a Usable Specification?

Let c be a CoT generated from the observation history o , and let $m(c)$ represent the action that c intends to perform, as described by it. We treat c as reliable only when it satisfies three criteria:

$$\text{Reliable}(c, o) = G(c, o) \wedge S(m(c), o) \wedge L(c, o). \quad (1)$$

We view a CoT as a candidate local specification: it identifies scene preconditions and uses them to justify an action whose execution should satisfy safety postconditions. Our annotation protocols aim to assess CoT reliability across three criteria:

Scene grounding (G). Every scene precondition used to justify the action must be supported by the model’s visual input. For example, “stop because the light is red” is grounded only when the relevant signal is visibly red. We also reject claims for which the input provides insufficient evidence, such as asserting that a lead vehicle is slower when relative motion cannot be established.

Action safety (S). The action prescribed in the CoT must satisfy the relevant safety postconditions: it must not create an observable hazard or violate an active traffic constraint if executed. Thus, stopping for a pedestrian in a crosswalk is safe, whereas nudging into an opposing lane to overtake is unsafe when an oncoming vehicle is approaching. This criterion assesses the action stated by the CoT, not whether the decoded trajectory actually realizes it.

Logical coherence (L). The grounded preconditions must plausibly imply the stated action and its intended postcondition. A slower lead vehicle together with a safe target-lane gap can justify a lane change. In contrast, changing into a lane occupied by a merging vehicle does not follow from the stated goal of yielding to that vehicle.

A CoT is a candidate specification only if all three criteria hold. Failures are non-exclusive: one CoT may, for example, hallucinate a lane and prescribe an unsafe maneuver into it.

B. Benchmark and Evaluation Protocol

Collection. Alpamayo 1.5 is an open 10B-parameter reasoning VLA that emits a natural-language CoT followed by a 6.4 s




Fig. 3. CoT reliability conditioned on each decision or causal-factor label; sample size is shown beside every label, and the dashed line marks overall reliability (66.7%). These are overlapping views of the same 150 pairs, not three disjoint partitions. A pair has at most one label in each decision channel (40 pairs state no longitudinal decision) but may have multiple causal factors.

ego trajectory [7], [9]. We run it in the released AlpaSim configuration [24] on NuRec reconstructed scenes, recording each emitted CoT and trajectory as one pair. Scenes carry metadata for layout, road type, weather, and traffic density. For coverage tracking and subgroup analysis, we label each CoT with GPT-5.5 using the AlpaMayo-R1 longitudinal-decision, lateral-decision, and causal-factor taxonomy [7]; these labels are not used as reliability, consistency, or safety ground truth. We first select scenes to cover the 40 provided scene-label categories, then favor rollouts with previously unseen taxonomy labels, and sample randomly after coverage saturates. The resulting DriveAlignBench includes 150 pairs covering all longitudinal decision categories and five lateral decision categories; out-of-lane nudge, pull-over/curb approach, and lateral-maneuver abort are absent because AlpaMayo rarely or never produces CoTs that GPT-5.5 maps to them.

Annotation. Two authors apply the rubric while viewing the CoT and the current and preceding camera frames supplied to the model. The trajectory is withheld so that CoT reliability is judged independently of CoT–trajectory consistency. Both annotators first label a 20-pair calibration set and refine the worked examples; one annotator then labels the remaining pairs, while the second independently labels a 10-pair subset. On this subset, the annotators agree on the final reliable/unreliable label for 9/10 pairs. For each pair, we retain the three reliability-criterion labels, a short justification for each violation, and the decision and causal-factor tags. Shared threats to validity are discussed in Section V-D.

C. Reliability Findings

Overall reliability (RQ1). 100/150 CoTs (66.7%) satisfy all three criteria. Of the 50 unreliable CoTs, 43 (86.0%) contain a grounding failure, 7 (14.0%) prescribe an unsafe action, and 6 (12.0%) are logically incoherent; 6 fail more than one criterion. Grounding is therefore the dominant obstacle to treating a CoT as a specification. Because sampling targets scene and behavior coverage, the rate characterizes this benchmark rather than the natural driving distribution. Fig. 2 shows this distinction in four pairs: two reliable CoTs drawn from very different scenes, and two unreliable CoTs that fail for different reasons—one a grounding error, the other an unsafe plan.

Which CoTs are reliable (RQ2). Fig. 3 shows the benchmark’s driving-decision/causal factor-conditioned views. Set-

speed tracking (8/10), lead following (40/52), and lane keeping (61/75) are the most reliable well-represented decision categories. Reliability falls sharply for lane changes (16/42), gap-searching (6/13), yielding (3/8), and acceleration for overtaking (3/10). Among causal factors, routing intent (11/25) and lane/laneline reasoning (44/82) are least reliable.

Overall, CoTs are more reliable for routine, directly observable driving behavior, and less reliable when they must infer lane topology, routing intent, or whether a spatial or temporal gap is available. For these weaker categories, AlpaMayo’s CoTs require additional reliability screening before being used as specifications.

IV. MONITORING FOR INCONSISTENCY DETECTION

The reliability study in Section III shows that although not all generated CoTs satisfy our reliability criteria, a substantial subset can serve as trustworthy specifications. For these CoTs, the remaining runtime verification task is to check whether the generated trajectory fulfills the expressed intent.

We therefore study monitors that compare a high-level natural-language specification against a low-level geometric plan. Given a reliable CoT and a planned trajectory, the monitor should flag cases where the plan contradicts, omits, or materially deviates from the maneuver stated in the CoT.

A. Problem Formulation

Challenges. Monitoring CoT–trajectory consistency is difficult because the two artifacts describe driving behavior at different levels of abstraction. The CoT states intent in natural language, while the trajectory records future ego motion as geometry and kinematics. A monitor must bridge this gap along two dimensions.

(C1) *Open-vocabulary ego-maneuver extraction.* The CoT c is a free-form natural-language artifact rather than a closed-set action label. The same ego maneuver may be expressed in many ways, such as *move over*, *merge right*, or *position for the upcoming turn*. A monitor must therefore recover the intended ego maneuver from open-vocabulary CoTs.

(C2) *Semantic gap between CoT and trajectory.* The maneuvers stated in c are often scene-relative; they reference lanes, map elements, and other agents, e.g., *keep lane*, *stop before the crosswalk*, or *move left to avoid an obstacle*. The

trajectory τ , by contrast, is typically a sequence of waypoints that yield speed, acceleration, and heading rate but not the entities c refers to. This semantic gap means that whether a stated maneuver is realized often cannot be decided from ego motion alone: *keep lane* requires the lane geometry, *stop before the crosswalk* the crosswalk’s position, and *move left to avoid an obstacle* the obstacle’s location. Therefore, deciding consistency may require additional information, such as map context.

We formalize this comparison as follows. A consistency monitor M takes a CoT c and a representation of the planned trajectory produced by a transform ϕ , and returns

$$\hat{y} = M(c, \phi(\tau, \xi)), \hat{y} \in \{\text{CONSISTENT}, \text{INCONSISTENT}\}.$$

Here, τ denotes the planned trajectory, and $\xi \in \Xi$ denotes optional non-trajectory context used to construct the monitor’s evidence representation. When the monitor uses only τ , we set $\xi = \emptyset$ and write $\phi(\tau)$ for $\phi(\tau, \emptyset)$. Consistency refers to alignment between the planned ego behavior and the maneuver stated in the CoT. For Alpamayo, a pair is inconsistent if the generated trajectory contradicts, omits, or fundamentally deviates from the stated maneuver. Examples include a trajectory that keeps the lane when the CoT specifies a lane change, merely decelerates when the CoT calls for a complete stop, or accelerates through a conflict zone when the CoT states a yield.

B. Monitoring Approaches

Our central design question is which trajectory representation makes c and τ comparable. We therefore design three monitors that progressively enrich the evidence available for consistency checking, as summarized in Fig. 4. The rule-based monitor maps both artifacts to a shared Alpamayo action vocabulary, giving an interpretable baseline based on discrete label matching. The LLM monitor keeps the trajectory in an ego-frame waypoint representation but replaces fixed phrase matching with an LLM judge, improving open-vocabulary maneuver extraction. The F-LLM monitor keeps the same judge interface and further adds lane-relative trajectory evidence, explicitly targeting the representation gap for behaviors such as lane keeping, nudging, and lane changes.

1) *Rule-Based Label-Matching Monitor*: The rule-based monitor maps both artifacts into a shared closed-set vocabulary of ego behaviors and then performs exact label matching. A CoT parser P_c and a trajectory parser P_τ produce action sequences

$$P_c(c) \rightarrow \mathbf{a}_c \in A^*, \quad P_\tau(\tau) \rightarrow \mathbf{a}_\tau \in A^*,$$

where $\mathbf{a}_c$ contains the maneuvers stated in the CoT and $\mathbf{a}_\tau$ contains trajectory-derived behaviors. The monitor returns

$$M_{\text{rule}}(c, \phi_{\text{act}}(\tau)) = \text{Match}(\mathbf{a}_c, \mathbf{a}_\tau),$$

where $\phi_{\text{act}}(\tau)$ is the trajectory’s action-label evidence. Match returns CONSISTENT only when every maneuver stated in the CoT appears, in order, as the same trajectory action label; otherwise it returns INCONSISTENT. For example, a CoT that

states *decelerate* then *stop* is consistent only if both labels appear in the trajectory-derived sequence.

Alpamayo instantiation. We use Alpamayo-R1’s longitudinal and lateral meta-action vocabulary as A , matching its RL post-training alignment between generated reasoning and trajectory-derived ego behavior. The longitudinal channel covers speed behavior such as *strong_accelerate*, *gentle_decelerate*, *stop*, and *maintain_speed*; the lateral channel covers steering behavior such as *sharp_steer_left*, *steer_right*, and *go_straight*.

We instantiate P_c as a dictionary-based phrase-to-action parser, with a total of 200 phrases mined from a development split disjoint from the evaluation set and augmented with paraphrases. Given a CoT, P_c maps matched maneuver phrases to longitudinal or lateral action labels in the shared Alpamayo vocabulary. We instantiate P_τ as a kinematic parser that thresholds per-segment speed, acceleration, and curvature into the same action channels. The full vocabulary, phrase dictionary, thresholds, matcher, and implementation details are provided in the public repository.

2) *LLM Monitor*: The LLM monitor replaces closed-set phrase matching with a prompted LLM judge over the CoT and ego-frame trajectory evidence. The complete monitor constructs the evidence, queries the judge, and thresholds its alignment score; the judge also returns a short rationale for error analysis. We use a fixed prompt with a coordinate-frame legend, few-shot examples, and a scoring rubric. This addresses the open-vocabulary nature of CoTs: instead of requiring every maneuver phrase to map to a predefined action label, the judge can interpret paraphrases such as “move over,” “merge right,” or “position for the turn” against the same ego-motion evidence used by the rule-based baseline.

$$(r, s) = \text{LLM}(\text{prompt}(c, \phi_{\text{ego}}(\tau)))$$

$$M_{\text{LLM}}(c, \phi_{\text{ego}}(\tau)) = \begin{cases} \text{INCONSISTENT} & \text{if } s \leq \theta \\ \text{CONSISTENT} & \text{otherwise} \end{cases}$$

Alpamayo instantiation. We set $\theta = 2$, so scores $s \leq 2$ are INCONSISTENT. To form $\phi_{\text{ego}}(\tau)$, we derive ego-motion quantities from the original 10 Hz trajectory and downsample to approximately 1.25 Hz by retaining one waypoint every 0.8 s. This coarser resolution keeps the maneuver-level shape needed for consistency checking while reducing prompt length and avoiding over-emphasis on low-level control smoothness. Each retained timestep contains ego position (x, y) , speed v , lateral velocity $\dot{y}$, and longitudinal acceleration a ; the prompt also includes duration, path length, final longitudinal and lateral displacement, lateral-deviation range, and mean and maximum speed. Prompts and implementation details are provided in the public repository.

3) *F-LLM Monitor: Enriching the Frame of Reference*: The monitor $M_{\text{F-LLM}}$ keeps the LLM judge unchanged but augments $\phi_{\text{ego}}(\tau)$ with lane-referenced geometry. This targets




Fig. 4. Progression of the three consistency monitors. The rule-based monitor compares discrete action sequences in a shared vocabulary. The LLM monitor replaces fixed phrase matching with an LLM judge over the CoT and a serialized ego-frame trajectory. F-LLM keeps the same judge interface but enriches the trajectory evidence with lane-relative quantities, making lane-centered CoT intents directly comparable to the planned trajectory.

CoT actions defined in the road frame: “keep lane,” “nudge left,” and “change lanes to the right” are better measured by offset from the lane center than by ego-frame lateral displacement. Without this frame, ordinary curve following can appear as a lateral maneuver under a fixed initial ego-frame. It can also mask a semantically important pedestrian-clearance nudge when road curvature in the opposite direction dominates the ego-frame lateral displacement. Fig. 5 shows a curved-road case where raw ego-frame motion overstates lateral departure, while lane-relative evidence correctly indicates lane keeping.

Starting from a high-definition map¹, we retrieve connected route candidates from the ego pose, select the route whose projection on τ minimizes the sum of squared lateral offsets, and compute $\phi_{\text{lane}}(\tau, \text{route})$ by comparing τ against that reference. The resulting offsets expose whether the ego stays near the lane center, drifts, nudges, or changes lanes.

Alpamayo instantiation. $M_{\text{F-LLM}}$ queries AlpaSim for the local road map, builds a directed lane graph from lane-center polylines, and enumerates route candidates from the ego’s initial pose. It selects $\text{route} = \arg \min_{\text{route}_j} \sum_{t=1}^T e_{t,j}^2$, where $e_{t,j}$ is the signed lateral offset of trajectory point t from candidate route r_j , then projects all trajectory points onto this fixed reference. The F-LLM prompt contains the raw ego-frame evidence plus longitudinal progress l_t , signed lateral offset e_t , offset change $\Delta e_t = e_t - e_1$, and lateral velocity $\dot{e}_t$ for lane-centered judgments. Prompts and implementation details are provided in the public repository.

Taken together, the three monitors separate the two challenges identified at the start of the section. The rule-based monitor establishes a transparent closed-vocabulary baseline,

¹We assume that access to high definition maps, such as those used extensively by autonomous vehicles as part of their driving stack is available [25].





Fig. 5. Frame-of-reference effect for a lane-keeping example (CoT: “keep lane to continue since the lane is clear ahead”). The trajectory τ follows a left-curving lane. (a) In the raw ego-frame representation $\phi_{\text{ego}}(\tau)$, the same lane-following trajectory is measured against a straight ray fixed to the ego’s initial heading, producing a large apparent lateral displacement ($e_t = +35.2$ m). (b) In the lane-relative representation $\phi_{\text{lane}}(\tau, \text{route})$, τ is measured against the routed lane center, and the signed offset remains within $+3.7$ m, more closely representing the behavior as lane-keeping.

but is limited by fixed phrase extraction. The LLM monitor relaxes that restriction by interpreting open-vocabulary CoTs, but still judges them against raw ego-frame motion. F-LLM then addresses the remaining representation gap for lane-centered maneuvers by adding route-relative trajectory evidence. This progression sets up the evaluation in Section V, where we measure how each additional representation choice changes inconsistency detection accuracy.



(a) Consistent & safe

“Change lane to the left due to slower traffic ahead and a clear gap.”

Planned path (orange) shifts one lane left, as stated, and stays clear of other traffic.



(b) Inconsistent & unsafe

“Change lane to the left due to slower vehicles ahead...”

Path instead swings left across the intersection into oncoming traffic, not a lane change.

Fig. 6. Representative CONSISTENT/INCONSISTENT CoT-trajectory pairs from the reliable subset with independent trajectory-safety labels. Each panel pairs the front-camera view (left) with the planned trajectory τ in bird’s-eye view (right; ego box in green, planned path in orange, lanes in blue). (a) Consistent and safe: τ realizes the stated left lane change and stays clear of traffic. (b) Inconsistent and unsafe: the CoT states a left lane change, but τ swings left across the intersection into oncoming traffic. Consistency cannot be read from the CoT alone; it requires checking the trajectory against the scene, motivating the lane-relative monitor F-LLM. Panel (b) illustrates the inconsistency–safety link quantified in RQ4 (Table II); the two labels are independently annotated and need not coincide.

V. EVALUATING THE MONITORING APPROACHES

This section provides an evaluation of the monitor approaches of Section IV and studies the safety relevance of CoT-trajectory inconsistency in an open-loop setting.

A. Research Questions

We organize the evaluation around two questions.

- **RQ3:** How accurately do the monitoring approaches distinguish CONSISTENT from INCONSISTENT CoT-trajectory pairs on reliable CoTs, and how does performance vary with the underlying model and the trajectory representation?
- **RQ4:** Is CoT-trajectory inconsistency associated with trajectory safety, and how well does it function as a filter for unsafe trajectories?

B. Experimental Setup

a) Objects of Evaluation: We restrict the evaluation to the reliable subset of 100 pairs from the benchmark in Section III so that accuracy is not confounded by ungrounded, unsafe-to-follow, or incoherent CoTs whose consistency cannot be meaningfully checked. Following the protocol in Section III, we annotate each pair with two additional binary labels: a CoT-trajectory consistency label and a trajectory-safety label. The *consistency* label records whether τ realizes the maneuver asserted in c , in the sense defined in Section IV. Each annotator is given the CoT text c , the trajectory τ (including its BEV rendering, raw waypoints, and kinematic summary), the relevant ego-camera frame, and the full video and metadata for context. Fig. 6 shows annotated examples of a consistent case (a) and an inconsistent case (b). The *safety* label, assigned independently of c , marks whether the realized trajectory τ would create a hazard or a traffic violation; for example, the trajectory in Fig. 6(b) abandons the stated lane change and swings left across the intersection into oncoming traffic, and is labeled unsafe. Of the 100 pairs, 26 are INCONSISTENT and 74 CONSISTENT, and 9 are UNSAFE and 91 SAFE. Agreement among annotators on the doubly annotated subset is 10/10 for both consistency and safety labels.

b) Treatments: We evaluate the three monitors of Section IV:

- **Rule-based:** a deterministic parse of the meta-action stated in c against the meta-action implied by τ .
- **LLM:** the judge LLM sees c and the raw ego-frame waypoint serialization $\phi_{\text{ego}}(\tau)$.
- **F-LLM:** the judge LLM sees c , the raw ego-frame serialization $\phi_{\text{ego}}(\tau)$, and the lane-relative representation $\phi_{\text{lane}}(\tau, \text{route})$.

The monitor families differ in their required artifacts: the rule-based monitor uses a phrase–action dictionary and hand-set trajectory thresholds; LLM uses a judge model and fixed scoring prompt over raw waypoints; and F-LLM adds a routed lane graph plus lane-relative trajectory projection. The LLM and F-LLM monitors use a single judge LLM to score each pair. We instantiate them with three diverse judge models:

- **Qwen3.5-4B-FP8**, a small, fast open-weight model served locally from the FP8-quantized model card [26]. Queried in thinking mode;
- **Kimi K2.5**, an open-weight multimodal mixture-of-experts model with 1T total and 32B activated parameters [27]. Queried in thinking mode;
- **GPT-5.5**, a proprietary frontier reasoning model, queried with reasoning effort `high` [28].


To minimize nondeterminism, all models receive identical inputs and a fixed decoding configuration (`seed = 42`). Because the hosted Kimi K2.5 and GPT-5.5 still exhibit residual serving-side nondeterminism, we repeat each hosted configuration three times and report the representative median-F1 run in Table I; run-to-run F1 ranges remain modest (at most 0.04 for Kimi K2.5 and 0.05 for GPT-5.5) relative to the trajectory-representation gaps analyzed below.

c) Metrics: Treating INCONSISTENT as the positive class, we report precision, recall, and F1. For RQ4, we cross-tabulate CoT-trajectory consistency labels with trajectory-safety labels, and evaluate whether inconsistency can act as a safety filter by measuring its recall and precision for unsafe trajectories.



(a) Scene + chain-of-thought

“Turn right to avoid the cones blocking the left fork.”




(b) Raw ego-frame reference

$\phi_{\text{ego}}(\tau)$

raw-waypoint monitor: 5/5 CONSISTENT

(wrong)



(c) Lane-relative reference

$\phi_{\text{lane}}(\tau, \text{route})$

F-LLM: 1/5 INCONSISTENT (correct)

Fig. 7. Example of an inconsistent CoT-trajectory pair where a richer frame of reference lets F-LLM overcome a raw-waypoint LLM-monitor error. The CoT in (a) says to *turn right* to avoid cones blocking the left fork, but the realized trajectory τ does not complete the right turn and instead turns left soon afterward. In (b), using the raw ego-frame reference $\phi_{\text{ego}}(\tau)$, the raw-waypoint monitor’s judge LLM assigns the pair a maximal consistency score (5/5), because τ shows predominantly rightward motion relative to a straight ray fixed to the ego’s initial heading. This is a false negative. In (c), using the lane-relative reference $\phi_{\text{lane}}(\tau, \text{route})$, F-LLM correctly labels the pair INCONSISTENT (1/5): relative to the routed lane, which bends right, the same trajectory drifts left of the lane center rather than following the intended turn. The lane-relative view exposes the divergence that the raw ego frame conflates with road geometry.

TABLE I

COT-TRAJECTORY INCONSISTENCY DETECTION ON THE RELIABLE SUBSET (100: 26 INCONSISTENT, 74 CONSISTENT). INCONSISTENT IS THE POSITIVE CLASS. THE FOUR COUNT COLUMNS FORM PER-ROW CONFUSION MATRICES OVER MATCHED ENTRIES. THE KIMI K2.5 AND GPT-5.5 ROWS SHOW THE REPRESENTATIVE MEDIAN-F1 RUN ACROSS THREE REPEATS. THE RULE-BASED AND QWEN3.5-4B METHODS USE A SINGLE RUN BECAUSE THEIR DETERMINISM CAN BE CONTROLLED. QWEN3.5-4B REFERS TO QWEN3.5-4B-FP8. METRICS ARE ROUNDED TO THE NEAREST 2 DIGITS. BEST EXACT F1 SCORE IS **BOLD**.

Monitor	Judge LLM	Confusion (counts)				Detection metrics		
		TP	FP	FN	TN	Prec.	Rec.	F1
Rule-based	N/A	9	13	17	61	0.41	0.35	0.38
LLM	Qwen3.5-4B	23	33	3	41	0.41	0.88	0.56
	Kimi K2.5	22	23	4	51	0.49	0.85	0.62
	GPT-5.5	21	24	5	50	0.47	0.81	0.59
F-LLM	Qwen3.5-4B	26	44	0	30	0.37	1.00	0.54
	Kimi K2.5	21	9	5	65	0.70	0.81	0.75
	GPT-5.5	20	7	6	67	0.74	0.77	0.75

C. Results

a) *Detection accuracy (RQ3)*: Table I reports metrics and confusion matrices for the three monitoring approaches. F-LLM (GPT-5.5) performs best ($F_1 = 0.75$, precision 0.74, recall 0.77), a +0.13 absolute F_1 gain over the strongest raw-waypoint baseline and well above the rule-based monitor. The rule-based baseline has $F_1 = 0.38$, detecting 9/26 inconsistencies and producing 13 false positives.

Having stronger hosted monitors does not necessarily improve detection accuracy. For raw-waypoint baseline, Kimi K2.5 outperforms GPT-5.5 by 0.03 F_1 , while GPT-5.5 marginally outperforms Kimi K2.5 with lane-relative features. Adding lane-relative features raises F_1 by +0.16 for GPT-5.5

and +0.13 for Kimi K2.5. Kimi K2.5 with lane-relative evidence also outperforms GPT-5.5 with raw waypoints (0.75 vs. 0.59), further indicating that structured scene-relative context is more important than scaling the judge LLM.

Qwen3.5-4B-FP8 monitor exhibits a high-recall, low-precision profile distinct from the larger hosted judges. It achieves the highest recall across all configurations, identifying 23/26 raw waypoint inconsistencies (recall 0.88) and a perfect 26/26 under F-LLM (recall 1.00), outperforming both hosted judges. However, frequent false alarms yield low precision (0.41 and 0.37, respectively), causing its F_1 scores (0.56 and 0.54) to trail the hosted models.

For Qwen, lane-relative features offer no F_1 gain (0.56 $\rightarrow$ 0.54), but they push the monitor to perfect recall (1.00), which is valuable for safety-oriented detection. The tradeoff is that Qwen does not convert the added scene-relative evidence into higher precision, suggesting that the utility of scene-relative evidence is inherently judge-dependent.

F-LLM’s main gain is to significantly reduce false positives: FP count drop from 23 to 9 for Kimi K2.5 and from 24 to 7 for GPT-5.5. Corrected cases show that raw ego-frame waypoints often conflate road geometry with ego intent, since curve-following, lane divergence, or minor clearance adjustments can mimic large lateral maneuvers. Lane-relative evidence resolves this by capturing whether a trajectory remains within the reference lane, shifts marginally, or executes a sustained lane change. It recovers cases where ego-frame motion superficially matches the CoT but contradicts route geometry (e.g., Fig. 7, where a “turn right” rationale accompanies a left-turning trajectory).

The remaining LLM-monitor errors shared by raw-waypoint and lane-relative runs, point to three limitations of the current evidence representation and judging setup. First, some errors




Fig. 8. Shared evaluation error of the raw-waypoint and lane-relative monitors. The CoT asserts *keeping distance* to a slowly moving lead vehicle, and human annotators label the pair **CONSISTENT**; both LLM-based monitors instead score it **INCONSISTENT** (2/5), reading mild ego acceleration as gap-closing. The failure arises because “keep distance” is relational over the ego-lead gap, but the monitor evidence contains only ego motion, so agent-relative predicates require agent-relative traces.

arise from overly local word matching, where the judge may over-weight an action word in the CoT and ignore small trajectory variations that are compatible with the maneuver in context. For example, it may treat an “accelerate” rationale as inconsistent with a trajectory that briefly decelerates before continuing the intended maneuver. Second, some errors arise because the monitor evidence is still incomplete near maneuver boundaries. Distinguishing a large in-lane nudge from a lane change, or a turn from following a curved road, can require structured scene features beyond the ego trajectory and its offset from a selected reference route, including lane width and route intent. When these features are absent or only implicit in the supplied evidence, the judge may draw a different maneuver boundary from the annotator. Third, some CoT claims require scene relations that are absent from the monitor evidence. For example, “keep distance” is a temporal commitment over the ego-lead gap, not merely a command to slow down. Because neither raw waypoints nor F-LLM’s lane-relative evidence identifies the lead vehicle or tracks that gap, the judge must infer compliance indirectly from ego speed and acceleration, treating slowing as evidence of keeping distance and mild acceleration as possible gap closing. This explains shared failures such as Fig. 8. Richer trajectory frames reduce geometric confusions, but they do not by themselves resolve errors caused by surface wording, ambiguous maneuver boundaries, or missing agent-relative traces.

b) Safety relevance (RQ4): Using human consistency labels, we first ask whether inconsistency itself is safety-relevant. Table II cross-tabulates those labels with independent trajectory-safety annotations on the reliable subset. Of the 9 unsafe trajectories, 7 are **INCONSISTENT**, so inconsistency labels cover $7/9 = 77.8\%$ of unsafe trajectories; Fig. 6(b) illustrates one such case, where ego’s planned trajectory diverges from its stated maneuver in an unsafe way. Among inconsistent pairs, $7/26 = 26.9\%$ are unsafe, higher than $2/74 = 2.7\%$ among consistent pairs. Thus, inconsistency

TABLE II
RELATIONSHIP BETWEEN CoT-TRAJECTORY CONSISTENCY AND TRAJECTORY SAFETY ON THE RELIABLE SUBSET. TREATING UNSAFE AS POSITIVE AND INCONSISTENT AS THE POSITIVE PREDICTION, INCONSISTENCY RECOVERS 7/9 UNSAFE TRAJECTORIES. UNSAFE TRAJECTORIES ARE MORE COMMON AMONG INCONSISTENT PAIRS ($7/26 = 26.9\%$) THAN AMONG CONSISTENT PAIRS ($2/74 = 2.7\%$).

CoT-trajectory label	Unsafe	Safe	Total	Unsafe rate
INCONSISTENT	7	19	26	26.9%
CONSISTENT	2	72	74	2.7%
Total	9	91	100	9.0%
Signal				Prec. Rec. F1
INCONSISTENT $\Rightarrow$ UNSAFE				0.27 0.78 0.40

is safety-relevant and can add a monitoring layer, but is insufficient as the sole safety-monitoring interface.

Across the deployed F-LLM monitors, all runs classify the 7 unsafe, human-labeled **INCONSISTENT** CoT-trajectory pairs as **INCONSISTENT**. Raw-waypoint LLM monitors recovered fewer of these cases across complete hosted repeats (4/7–6/7). This suggests that a performant consistency monitor can help identify safety risk. The remaining 2 unsafe trajectories are **CONSISTENT** with the CoT and therefore outside the scope of a faithful consistency monitor.

c) Monitor latency: Inference latency remains a practical limitation. Used off-the-shelf, these monitors are too slow for online runtime deployment: hosted GPT-5.5 and Kimi K2.5 require several seconds to tens of seconds per query, and their proprietary APIs prevent context-handling or serving-stack optimization. The local Qwen3.5-4B-FP8 monitor achieves sub-second inference on a single NVIDIA A100-80GB GPU, but its high false-positive rate limits its usefulness as a standalone runtime monitor. We leave optimizations such as distilling larger models into smaller, faster student models for monitoring tasks to future work.

D. Threats to Validity

External validity. Our benchmark uses one reasoning VLA (Alpamayo 1.5), one simulator/reconstruction pipeline, and a diversity-oriented sample rather than a naturalistic driving input distribution. VLA reasoning models for AVs and their traces are still emerging, so future models may expose different CoT formats and failure modes. Simulation also abstracts real-world perception noise, vehicle dynamics, and operational-design-domain constraints; reported rates should therefore be read as conditional on this setup rather than as population estimates for deployed AVs.

Internal validity. Outcomes may be affected by implementation choices in trajectory parsing, route selection, lane projection, prompt construction, and result aggregation. We reduce avoidable variation by using identical inputs, fixed decoding settings, and repeated hosted-model runs, but residual implementation or serving-side effects may remain. Annotation is another source of subjectivity: labels follow an explicit rubric with worked examples, yet most pairs receive

a single final label and the doubly annotated subset shows 9/10 for reliability and 10/10 agreement for both consistency and safety. The chosen judge models, prompts, and sampled frames may also affect monitor effectiveness. Finally, the rule-based monitor’s thresholds, phrase dictionary, and matcher are our reconstruction of Alpamayo-R1’s action-consistency procedure from the available description in [7]; they may differ from the original implementation.

Construct validity. Reliability, consistency, and safety are operational labels for broader concepts. Our rubrics are designed to make these judgments reproducible, but alternative rubrics could handle ambiguous CoTs, defensive driving, near misses, or traffic-law violations differently. Restricting RQ3 and RQ4 to reliable CoTs avoids confounding monitor accuracy with unreliable specifications, but it also conditions the conclusions on cases where the CoT can be treated as a meaningful local specification. LLM-monitor outputs are further reduced from graded scores to binary decisions by a fixed threshold, so other operating points could trade precision for recall.

Conclusion validity. The evaluation contains 100 reliable pairs and only 9 unsafe trajectories, so subgroup rates and the inconsistency–safety association are indicative rather than definitive. We therefore report descriptive counts, precision, recall, and F1 instead of claiming statistical significance or causal effects. Larger randomly sampled datasets, closed-loop evaluation, and additional VLA/judge-model combinations are needed to confirm the observed trends.

VI. CONCLUSION

This paper studies whether chain-of-thought (CoT) traces from driving VLAs can serve as specifications for runtime V&V. We identify two requirements for this idea to be meaningful: the CoT must first be a reliable candidate specification, and the generated trajectory must then be comparable to that specification. On DriveAlignBench, a 150-pair Alpamayo 1.5 benchmark, 100/150 CoTs satisfy our grounding, safety, and coherence criteria. Reliability is strongest for routine, directly observable behavior such as lead following, lane keeping, and set-speed tracking, and weaker for decisions that require inferring lane topology, routing intent, or spatial and temporal gaps. These results suggest that CoTs can provide useful runtime specifications, but only selectively and with explicit reliability filtering.

We then develop rule-based, raw-waypoint LLM, and lane-relative F-LLM monitors for checking whether a trajectory realizes a reliable CoT. On the reliable subset, the best F-LLM monitor reaches $F_1 = 0.75$ and improves over both the strongest raw-waypoint LLM baseline and the rule-based monitor, showing the value of exposing scene-relative trajectory evidence to the judge. CoT–trajectory inconsistency is also safety-relevant: 7/9 unsafe trajectories are INCONSISTENT, although inconsistency alone is not a complete safety monitor. The remaining failures point to broader limits of LLM-based monitoring: judges can over-weight surface action words, draw different boundaries between nearby maneuvers, and lack the agent-relative traces needed to check claims such as “keep

distance” from ego motion alone. Future work should broaden the evaluation across additional VLAs, emerging world-model-based driving architectures that expose comparable reasoning or plan traces, and real-world data, with deeper analysis of which scenario factors drive driving policy and monitor failures. More importantly, monitors should move beyond lane geometry to agent-relative and temporal traces, so claims such as “keep distance,” “yield,” or “merge behind” can be checked against tracked scene state. Closed-loop studies should then test whether these monitors can trigger fallback or replanning before inconsistent trajectories become hazards.

REFERENCES

- [1] National Transportation Safety Board, “Collision between vehicle controlled by developmental automated driving system and pedestrian, tempe, arizona, march 18, 2018,” National Transportation Safety Board, Washington, DC, Highway Accident Report NTSB/HAR-19/03, 2019, pB2019-101402. [Online]. Available: <https://www.nts.gov/investigations/AccidentReports/Reports/HAR1903.pdf>
- [2] —, “Collision between a sport utility vehicle operating with partial driving automation and a crash attenuator, mountain view, california, march 23, 2018,” National Transportation Safety Board, Washington, DC, Highway Accident Report NTSB/HAR-20/01, 2020, pB2020-100112. [Online]. Available: <https://www.nts.gov/investigations/AccidentReports/Reports/HAR2001.pdf>
- [3] National Highway Traffic Safety Administration, “NHTSA announces consent order with cruise after company failed to fully report crash involving pedestrian,” Press release, Sep. 2024, september 30, 2024. [Online]. Available: <https://www.nhtsa.gov/press-releases/consent-order-cruise-crash-reporting>
- [4] L. Chen, P. Wu, K. Chitta, B. Jaeger, A. Geiger, and H. Li, “End-to-end autonomous driving: Challenges and frontiers,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 46, no. 12, pp. 10 164–10 183, 2024. [Online]. Available: <https://arxiv.org/abs/2306.16927>
- [5] Wayve, “LINGO-2: Driving with natural language,” Wayve blog, Apr. 2024, published April 17, 2024. [Online]. Available: <https://wayve.ai/thinking/lingo-2-driving-with-language/>
- [6] J.-J. Hwang, R. Xu, H. Lin, W.-C. Hung, J. Ji, K. Choi, D. Huang, T. He, P. Covington, B. Sapp, Y. Zhou, J. Guo, D. Anguelov, and M. Tan, “EMMA: End-to-end multimodal model for autonomous driving,” *Transactions on Machine Learning Research*, 2025. [Online]. Available: <https://arxiv.org/abs/2410.23262>
- [7] NVIDIA, Y. Wang, W. Luo, J. Bai, Y. Cao, T. Che, K. Chen, Y. Chen, J. Diamond, Y. Ding, W. Ding, L. Feng, G. Heinrich, J. Huang, P. Karkus, B. Li, P. Li, T.-Y. Lin, D. Liu, M.-Y. Liu, L. Liu, Z. Liu, J. Lu, Y. Mao, P. Molchanov, L. Pavao, Z. Peng, M. Ranzinger, E. Schmerling, S. Shen, Y. Shi, S. Tariq, R. Tian, T. Wekel, X. Weng, T. Xiao, E. Yang, X. Yang, Y. You, X. Zeng, W. Zhang, B. Ivanovic, and M. Pavone, “Alpamayo-R1: Bridging reasoning and action prediction for generalizable autonomous driving in the long tail,” *arXiv preprint arXiv:2511.00088*, 2025. [Online]. Available: <https://arxiv.org/abs/2511.00088>
- [8] K. Renz, L. Chen, E. Arani, and O. Sinavski, “SimLingo: Vision-only closed-loop autonomous driving with language-action alignment,” in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025. [Online]. Available: <https://arxiv.org/abs/2503.09594>
- [9] NVIDIA, “Alpamayo 1.5 model card,” Hugging Face model repository, 2026, release date: March 19, 2026. Accessed: 2026-07-01. [Online]. Available: <https://huggingface.co/nvidia/Alpamayo-1.5-10B>
- [10] F. Toledo, T. Woodlief, S. Elbaum, and M. B. Dwyer, “Specifying and monitoring safe driving properties with scene graphs,” in *2024 IEEE International Conference on Robotics and Automation (ICRA)*. IEEE, 2024, pp. 15 577–15 584.
- [11] T. Woodlief, F. Toledo, M. B. Dwyer, and S. Elbaum, “Scene flow specifications: Encoding and monitoring rich temporal safety properties of autonomous systems,” *Proceedings of the ACM on Software Engineering*, vol. 2, no. FSE, pp. 2524–2547, 2025. [Online]. Available: <https://doi.org/10.1145/3729382>

- [12] T. Wang, E. Xie, R. Chu, Z. Li, and P. Luo, "DriveCoT: Integrating chain-of-thought reasoning with end-to-end driving," 2024. [Online]. Available: <https://arxiv.org/abs/2403.16996>
- [13] X. Tian, J. Gu, B. Li, Y. Liu, Y. Wang, Z. Zhao, K. Zhan, P. Jia, X. Lang, and H. Zhao, "DriveVLM: The convergence of autonomous driving and large vision-language models," *arXiv preprint arXiv:2402.12289*, 2024. [Online]. Available: <https://arxiv.org/abs/2402.12289>
- [14] J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia, E. Chi, Q. V. Le, and D. Zhou, "Chain-of-thought prompting elicits reasoning in large language models," in *Advances in Neural Information Processing Systems*, vol. 35, 2022, pp. 24 824–24 837. [Online]. Available: <https://arxiv.org/abs/2201.11903>
- [15] Y. Luo, F. Li, S. Xu, Z. Lai, L. Yang, Q. Chen, Z. Luo, Z. Xie, S. Jiang, J. Liu, L. Chen, B. Wang, and Z.-x. Yang, "AdaThinkDrive: Adaptive thinking via reinforcement learning for autonomous driving," *arXiv preprint arXiv:2509.13769*, 2025. [Online]. Available: <https://arxiv.org/abs/2509.13769>
- [16] Z. Yuan, C. Qian, J. Tang, R. Chen, Z. Song, L. Sun, X. Chu, Y. Cai, D. Zhang, and S. Li, "AutoDrive-R²: Incentivizing reasoning and self-reflection capacity for VLA model in autonomous driving," *arXiv preprint arXiv:2509.01944*, 2025. [Online]. Available: <https://arxiv.org/abs/2509.01944>
- [17] M. Turpin, J. Michael, E. Perez, and S. R. Bowman, "Language models don't always say what they think: Unfaithful explanations in chain-of-thought prompting," in *Advances in Neural Information Processing Systems*, vol. 36, 2023. [Online]. Available: <https://arxiv.org/abs/2305.04388>
- [18] T. Lanham, A. Chen, A. Radhakrishnan, B. Steiner, C. Denison *et al.*, "Measuring faithfulness in chain-of-thought reasoning," *arXiv preprint arXiv:2307.13702*, 2023. [Online]. Available: <https://arxiv.org/abs/2307.13702>
- [19] N. Mayumu, X. Deng, and P. Mukala, "Is VLA reasoning faithful? probing safety of chain-of-causation in autonomous driving models," *arXiv preprint arXiv:2605.17268*, 2026. [Online]. Available: <https://arxiv.org/abs/2605.17268>
- [20] Y. Wu, A. Li, T. Hermans, F. Ramos, A. Bajcsy, and C. Pérez-DÁrpino, "Do what you say: Steering vision-language-action models via runtime reasoning-action alignment verification," *arXiv preprint arXiv:2510.16281*, 2025. [Online]. Available: <https://arxiv.org/abs/2510.16281>
- [21] S. Shalev-Shwartz, S. Shammah, and A. Shashua, "On a formal model of safe and scalable self-driving cars," *arXiv preprint arXiv:1708.06374*, 2017. [Online]. Available: <https://arxiv.org/abs/1708.06374>
- [22] D. Hendrycks and K. Gimpel, "A baseline for detecting misclassified and out-of-distribution examples in neural networks," in *International Conference on Learning Representations (ICLR)*, 2017. [Online]. Available: <https://arxiv.org/abs/1610.02136>
- [23] C. B. Kuhn, M. Hofbauer, G. Petrovic, and E. Steinbach, "Introspective black box failure prediction for autonomous driving," in *2020 IEEE Intelligent Vehicles Symposium (IV)*, 2020, pp. 1907–1913.
- [24] NVIDIA, Y. Cao, R. de Lutio, S. Fidler, G. Garcia Cobo, Z. Gojcic, M. Igl, B. Ivanovic, P. Karkus, J. Martinez Esturo, M. Pavone, A. Smith, E. Tanimura, M. Tyszkiewicz, M. Watson, Q. Wu, and L. Zhang, "AlpaSim: A modular, lightweight, and data-driven research simulator for autonomous driving," Oct. 2025. [Online]. Available: <https://github.com/NVlabs/alpasim>
- [25] Q. Li, Y. Wang, Y. Wang, and H. Zhao, "HDMaNet: An online HD map construction and evaluation framework," in *2022 International Conference on Robotics and Automation (ICRA)*, 2022, pp. 4628–4634. [Online]. Available: <https://arxiv.org/abs/2107.06307>
- [26] Red Hat AI, "Qwen3.5-4B-FP8-dynamic model card," Hugging Face model repository, 2026, accessed: 2026-07-01. [Online]. Available: <https://huggingface.co/RedHatAI/Qwen3.5-4B-FP8-dynamic>
- [27] Kimi Team, "Kimi K2.5: Visual agentic intelligence," *arXiv preprint arXiv:2602.02276*, 2026. [Online]. Available: <https://arxiv.org/abs/2602.02276>
- [28] OpenAI, "Introducing GPT-5.5," OpenAI product release, Apr. 2026, published April 23, 2026. Accessed: 2026-07-01. [Online]. Available: <https://openai.com/index/introducing-gpt-5-5/>